

To make Medium work, we log user data. By using Medium, you agree to our [Privacy Policy](#), including cookie policy. ×

# Simple Reinforcement Learning with Tensorflow Part 0: Q-Learning with Tables and Neural Networks



Arthur Juliani [Follow](#)

Aug 25, 2016 · 6 min read



We'll be learning how to solve the OpenAI FrozenLake environment. Our version is a little less photo-realistic.

For this tutorial in my Reinforcement Learning series, we are going to be exploring a family of RL algorithms called Q-Learning algorithms. These are a little different than the policy-based algorithms that will be looked at in the the following tutorials (Parts 1–3). Instead of starting with a complex and unwieldy deep neural network, we will begin by implementing a simple lookup-table version of the algorithm, and then show how to implement a neural-network equivalent using Tensorflow. Given that we are going back to basics, it may be best to think of this as Part-0 of the series. It will hopefully give an intuition into what is really happening in Q-Learning that we can then build on going forward when we eventually combine the policy gradient and Q-learning approaches to build state-of-the-art RL agents (If you are more interested in Policy Networks, or already have a grasp on Q-Learning, feel free to start the tutorial series [here](#) instead).

Unlike policy gradient methods, which attempt to learn functions which directly map an observation to an action, Q-Learning attempts to learn the value of being in a given state, and taking a specific action there. While both approaches ultimately allow us to take intelligent actions given a situation, the means of getting to that action differ significantly. You may have heard about [DeepQ-Networks which can play Atari Games](#). These are really just larger and more complex implementations of the Q-Learning algorithm we are going to discuss here.

Top highlight

## Tabular Approaches for Tabular Environments

```
SFFF      (S: starting point, safe)
FHFH      (F: frozen surface, safe)
FFFH      (H: hole, fall to your doom)
```

To make Medium work, we log user data. By using Medium, you agree to our [Privacy Policy](#), including cookie policy. ×

If this post has been valuable to you, please consider [\*donating\*](#) to help support future tutorials, articles, and implementations. Any contribution is greatly appreciated!

If you'd like to follow my work on Deep Learning, AI, and Cognitive Science, follow me on Medium [@Arthur Juliani](#), or on Twitter [@awjliani](#).

. . .

*More from my Simple Reinforcement Learning with Tensorflow series:*

1. *[Part 0 — Q-Learning Agents](#)*
2. *[Part 1 — Two-Armed Bandit](#)*
3. *[Part 1.5 — Contextual Bandits](#)*
4. *[Part 2 — Policy-Based Agents](#)*
5. *[Part 3 — Model-Based RL](#)*
6. *[Part 4 — Deep Q-Networks and Beyond](#)*
7. *[Part 5 — Visualizing an Agent's Thoughts and Actions](#)*
8. *[Part 6 — Partial Observability and Deep Recurrent Q-Networks](#)*
9. *[Part 7 — Action-Selection Strategies for Exploration](#)*
10. *[Part 8 — Asynchronous Actor-Critic Agents \(A3C\)](#)*

Machine Learning

Artificial Intelligence

Neural Networks

Deep Learning

Robotics

### Discover Medium

Welcome to a place where words matter. On Medium, smart voices and original ideas take center stage - with no ads in sight. [Watch](#)

### Make Medium yours

Follow all the topics you care about, and we'll deliver the best stories for you to your homepage and inbox. [Explore](#)

### Become a member

Get unlimited access to the best stories on Medium — and support writers while you're at it. Just \$5/month. [Upgrade](#)

**Medium**

[About](#)

[Help](#)

[Legal](#)